

## Homework 1: Linear Regression and Neural Network Regression.

**Anay Abhijit Joshi**

**CS 5173: Deep Learning**  
**Level: Undergraduate**

### Step 1: Data

1) How many data samples are included in the dataset?

**3047**

```
[8] # Name: Anay Abhijit Joshi
import pandas as pd

# Let's specify the file path that contains the data-set
file_path = "cancer_reg.csv"
# This dataset contains cancer mortality rates information. The goal of this project/challenge
# is to predict the results of Cancer Mortality Rates. Therefore, the label is "TARGET_deathRate".
# I will load the dataset into a pandas data-frame, and handle special-character encoding
file = pd.read_csv(filepath_or_buffer=file_path, encoding="latin1")

[9] print(file.shape[0])
```

3047

2) Which problem will this dataset try to address?

**This dataset will predict the results of the Cancer Mortality Rates. For this problem, the “label” is “TARGET\_deathRate”.**

3) What is the minimum value and the maximum value in the dataset?

**minimum value = 0.0**

**maximum value = 10170292**

```
[10] print(file.min())
```

avgAnnCount	6.0
avgDeathsPerYear	3
TARGET_deathRate	59.7
incidenceRate	201.3
medIncome	22640
popEst2015	827
povertyPercent	3.2
studyPerCap	0.0
binnedInc	(34218.1, 37413.8]
MedianAge	22.3
MedianAgeMale	22.4
MedianAgeFemale	22.3
Geography	Abbeville County, South Carolina
AvgHouseholdSize	0.0221
PercentMarried	23.1
PctNoHS18_24	0.0
PctHS18_24	0.0
PctSomeCol18_24	7.1
PctBachDeg18_24	0.0
PctHS25_Over	7.5
PctBachDeg25_Over	2.5
PctEmployed16_Over	17.6
PctUnemployed16_Over	0.4
PctPrivateCoverage	22.3
PctPrivateCoverageAlone	15.7
PctEmpPrivCoverage	13.5
PctPublicCoverage	11.2
PctPublicCoverageAlone	2.6
PctWhite	10.199155
PctBlack	0.0
PctAsian	0.0
PctOtherRace	0.0
PctMarriedHouseholds	22.99249
BirthRate	0.0
dtype:	object

```
[12] float_or_int_only = file.select_dtypes([float, int])
print(float_or_int_only.min().min())
```

```
0.0
```

```
[17] print(file.max())
```

avgAnnCount	38150.0
avgDeathsPerYear	14010
TARGET_deathRate	362.8
incidenceRate	1206.9
medIncome	125635
popEst2015	10170292
povertyPercent	47.4
studyPerCap	9762.308998
binnedInc	[22640, 34218.1]
MedianAge	624.0
MedianAgeMale	64.7
MedianAgeFemale	65.7
Geography	Zavala County, Texas
AvgHouseholdSize	3.97
PercentMarried	72.5
PctNoHS18_24	64.1
PctHS18_24	72.5
PctSomeCol18_24	79.0
PctBachDeg18_24	51.8
PctHS25_Over	54.8
PctBachDeg25_Over	42.2
PctEmployed16_Over	80.1
PctUnemployed16_Over	29.4
PctPrivateCoverage	92.3
PctPrivateCoverageAlone	78.9
PctEmpPrivCoverage	70.7
PctPublicCoverage	65.1
PctPublicCoverageAlone	46.6
PctWhite	100.0
PctBlack	85.947799
PctAsian	42.619425
PctOtherRace	41.930251
PctMarriedHouseholds	78.075397
BirthRate	21.326165
dtype:	object

```
float_or_int_only = file.select_dtypes([float, int])
print(float_or_int_only.max().max())
```

```
10170292.0
```

4) How many features in each data samples?

**33 Features**

*Excludes 1 Label*

*Total Columns: 34*

```
✓ [29] # Rows and Columns  
0s print(file.shape)
```

```
⇒ (3047, 34)
```

```
✓ [30] # Total Columns  
0s print(file.shape[1])
```

```
⇒ 34
```

```
✓ [31] # Exclude 1 column of "LABEL" to get "FEATURES"  
0s print((file.shape[1]) - 1)
```

```
⇒ 33
```

5) Does the dataset have any missing information? E.g., missing features.

**Yes, the dataset has the missing information.**

**PctSomeCol18\_24 2285**

**PctEmployed16\_Over 152**

**PctPrivateCoverageAlone 609**

```
✓ [42] print(file.isnull().sum())  
0s
```

```
⇒ avgAnnCount      0  
   avgDeathsPerYear 0  
   TARGET_deathRate 0  
   incidenceRate    0  
   medIncome        0  
   popEst2015       0  
   povertyPercent   0  
   studyPerCap      0  
   binnedInc        0  
   MedianAge        0  
   MedianAgeMale    0  
   MedianAgeFemale  0  
   Geography        0  
   AvgHouseholdSize 0  
   PercentMarried   0  
   PctNoHS18_24     0  
   PctHS18_24       0  
   PctSomeCol18_24   2285  
   PctBachDeg18_24   0  
   PctHS25_Over      0  
   PctBachDeg25_Over 0  
   PctEmployed16_Over 152  
   PctUnemployed16_Over 0  
   PctPrivateCoverage 0  
   PctPrivateCoverageAlone 609  
   PctEmpPrivCoverage 0  
   PctPublicCoverage 0  
   PctPublicCoverageAlone 0  
   PctWhite         0  
   PctBlack         0  
   PctAsian         0  
   PctOtherRace     0  
   PctMarriedHouseholds 0  
   BirthRate        0  
   dtype: int64
```

```
✓ [43] print((file.isnull().sum())[file.isnull().sum() > 0])  
0s
```

```
⇒ PctSomeCol18_24      2285  
   PctEmployed16_Over  152  
   PctPrivateCoverageAlone 609  
   dtype: int64
```

6) What is the label of this dataset?

**TARGET\_deathRate**

7) How many percent of data will you use for **training, validation and testing**?

**80% - 10% - 10%**

8) What kind of data pre-processing will you use for your training dataset?

**In the given dataset, for data pre-processing, first, I filled the missing values in the columns with the column's data's "mean". This was done to make sure that no data is being lost. Moreover, the categorical features, such as "Geography" and "binnedInc" were label encoded using "LabelEncoder()" to convert these feature values into numerical data and then, use it. In addition, I also used "numpy's mathematical log function" for the features and the target variable ("TARGET\_deathRate"), for transformation of the skewed data to reduce skewness and make the data distribution more symmetric. Further, I also used "StandardScaler()" for standardizing the features by removing the mean and scaling to unit variance. Finally, I used this data by splitting it into 80%-20%-20% for training, validation, and testing, respectively, for a balanced model evaluation.**

## Step 2: Model

Model	Test R-squared	MSE
Linear regression	0.770306897015671	0.005662627419371372
DNN-16 ( $LR=0.001$ )	0.8526817893566186	0.0041490313299556375
DNN-30-8 ( $LR=0.001$ )	0.8152757268053347	0.005202525835337959
DNN-30-16-8 ( $LR=0.001$ )	0.8182953390336587	0.005117482270900689
DNN-30-16-8-4 ( $LR=0.001$ )	0.7761951143111193	0.006303181924793659

**Step 5: Model Selection****Epochs = 100****Batch Size = 64****R-Squared value**

Model	LR: 0.1 ( $R^2$ )	LR: 0.01 ( $R^2$ )	LR: 0.001 ( $R^2$ )	LR: 0.0001 ( $R^2$ )
Linear regression	-	-	-	-
DNN-16	0.7609437147162864	0.8462586399544709	0.8526817893566186	0.8567260630482327
DNN-30-8	0.6768140614561169	0.8189500334814785	0.8152757268053347	0.8150193659723532
DNN-30-16-8	0.6775849782102178	0.8090008832305213	0.8182953390336587	0.8157842936758392
DNN-30-16-8-4	0.7301531324384998	0.7875299336077706	0.7761951143111193	0.7762262288035106

**MSE (Mean Squared Error)**

Model	LR: 0.1 (MSE)	LR: 0.01 (MSE)	LR: 0.001 (MSE)	LR: 0.0001 (MSE)
Linear regression	-	-	-	-
DNN-16	0.006732718330837954	0.004329931219997128	0.0041490313299556375	0.004035129469621198
DNN-30-8	0.009102123753496269	0.005099043628700995	0.005202525835337959	0.005209745914397893
DNN-30-16-8	0.009080411856836763	0.005379248879072817	0.005117482270900689	0.005188202691784384
DNN-30-16-8-4	0.0075998961901140565	0.005983951055943129	0.006303181924793659	0.006302305624415157


## Step 6: Model Performance

As shown in the table in STEP 2, for **Linear Regression** model, the **MSE (Mean Squared Error)** value was **0.770306897015671** and **Test R-Squared** value was **0.005662627419371372**

## Linear Regression Model

Test R-Squared = 0.770306897015671

Mean Squared Error (MSE) = 0.005662627419371372



The best performing model, with respect to all the obtained values of R-Squared value and Mean Squared Error (MSE) value is as follows -

## DNN-16


Learning Rate (LR) = 0.0001

*Epochs = 100*

*Batch-Size = 64*

Mean Squared Error (MSE) = 0.004035129469621198

R-Squared = 0.8567260630482327



I set the number of training iterations over the dataset, or “EPOCHS” to “100”; and the “BATCH-SIZE” to “64” for evaluating the performance of all the models with different learning rates. Therefore, these were consistent across all the models...